



Bits, Bytes and Integers – Part 1

15-213/18-213/14-513/15-513: Introduction to Computer Systems
2nd Lecture, Aug. 30, 2018

Announcements

- Recitations are on Mondays, but next Monday (9/3) is Labor Day, so recitations are cancelled
- Linux Boot Camp Monday evening 7pm, Rashid Auditorium
- Lab 0 is now available via course web page and [Autolab](#).
 - Due Thu Sept. 6, 11:59pm
 - No grace days
 - No late submissions
 - Just do it!

Logistics


- Waitlist
 - 15-213: Mary Widom (marwidom@cs.cmu.edu)
 - 18-213: ECE Academic services
ece-asc@andrew.cmu.edu
 - 15-513: Mary Widom (marwidom@cs.cmu.edu)
 - 14-513: INI Enrollment (ini-enrollment@andrew.cmu.edu)
 - Please don't contact the instructors with waitlist questions.
- Autolab Accounts
 - Check whether you have one
 - If not, refer to Piazza @68

Today: Bits, Bytes, and Integers

- Representing information as bits
- Bit-level manipulations
- Integers
 - Representation: unsigned and signed
 - Conversion, casting
 - Expanding, truncating
 - Addition, negation, multiplication, shifting
 - Summary
- Representations in memory, pointers, strings

Everything is bits

- Each bit is 0 or 1
- By encoding/interpreting sets of bits in various ways
 - Computers determine what to do (instructions)
 - ... and represent and manipulate numbers, sets, strings, etc...
- Why bits? Electronic Implementation
 - Easy to store with bistable elements
 - Reliably transmitted on noisy and inaccurate wires



For example, can count in binary

- Base 2 Number Representation
 - Represent 15213_{10} as 11101101101101_2
 - Represent 1.20_{10} as $1.0011001100110011[0011]_{\dots 2}$
 - Represent 1.5213×10^4 as $1.1101101101101_2 \times 2^{13}$

Encoding Byte Values

- Byte = 8 bits
 - Binary 00000000_2 to 11111111_2
 - Decimal: 0_{10} to 255_{10}
 - Hexadecimal 00_{16} to FF_{16}
 - Base 16 number representation
 - Use characters '0' to '9' and 'A' to 'F'
 - Write $FA1D37B_{16}$ in C as
 - `0xFA1D37B`
 - `0xfa1d37b`

	Hex	Decimal	Binary
0	0	0000	
1	1	0001	
2	2	0010	
3	3	0011	
4	4	0100	
5	5	0101	
6	6	0110	
7	7	0111	
8	8	1000	
9	9	1001	
A	10	1010	
B	11	1011	
C	12	1100	
D	13	1101	
E	14	1110	
F	15	1111	

15213: 0011 1011 0110 1101
3 B 6 D

Example Data Representations

C Data Type	Typical 32-bit	Typical 64-bit	x86-64
char	1	1	1
short	2	2	2
int	4	4	4
long	4	8	8
float	4	4	4
double	8	8	8
pointer	4	8	8

Today: Bits, Bytes, and Integers

- Representing information as bits
- Bit-level manipulations
- Integers
 - Representation: unsigned and signed
 - Conversion, casting
 - Expanding, truncating
 - Addition, negation, multiplication, shifting
 - Summary
- Representations in memory, pointers, strings

Boolean Algebra

- Developed by George Boole in 19th Century
 - Algebraic representation of logic
 - Encode "True" as 1 and "False" as 0
- And
- Or
- Not
- Exclusive-Or (Xor)
- $\sim A = 1$ when $A=0$
- $A \wedge B = 1$ when both $A=1$ and $B=1$
- $A \vee B = 1$ when either $A=1$ or $B=1$
- $A \oplus B = 1$ when either $A=1$ or $B=1$, but not both

$\&$	0	1
0	0	0
1	0	1

\vee	0	1
0	0	1
1	1	1

\sim	
0	1
1	0

\wedge	0	1
0	0	1
1	1	0

General Boolean Algebras

- Operate on Bit Vectors
 - Operations applied bitwise
- $$\begin{array}{ccccccc} 01101001 & 01101001 & 01101001 & & 01101001 \\ \underline{\oplus} & \underline{\vee} & \underline{\wedge} & \underline{\sim} & \underline{\oplus} \\ 01000001 & 01111101 & 00111100 & & 10101010 \end{array}$$
- All of the Properties of Boolean Algebra Apply

Example: Representing & Manipulating Sets

- Representation
 - Width w bit vector represents subsets of $\{0, \dots, w-1\}$
 - $a_j = 1$ if $j \in A$
 - 01101001 {0, 3, 5, 6}
 - 76543210
 - 01010101 {0, 2, 4, 6}
 - 76543210
- Operations
 - $\&$ Intersection 01000001 {0, 6}
 - \vee Union 01111101 {0, 2, 3, 4, 5, 6}
 - \wedge Symmetric difference 00111100 {2, 3, 4, 5}
 - \sim Complement 10101010 {1, 3, 5, 7}

Bit-Level Operations in C

- Operations $\&$, \vee , \sim , \wedge Available in C
 - Apply to any "integral" data type
 - long, int, short, char, unsigned
 - View arguments as bit vectors
 - Arguments applied bit-wise
- Examples (Char data type)
 - $\sim 0x41 \rightarrow$
 - $\sim 0x00 \rightarrow$
 - $0x69 \& 0x55 \rightarrow$
 - $0x69 \vee 0x55 \rightarrow$

Hex	Decimal	Binary
0	0	0000
1	1	0001
2	2	0010
3	3	0011
4	4	0100
5	5	0101
6	6	0110
7	7	0111
8	8	1000
9	9	1001
A	10	1010
B	11	1011
C	12	1100
D	13	1101
E	14	1110
F	15	1111

Bit-Level Operations in C

- Operations $\&$, \vee , \sim , \wedge Available in C
 - Apply to any "integral" data type
 - long, int, short, char, unsigned
 - View arguments as bit vectors
 - Arguments applied bit-wise
- Examples (Char data type)
 - $\sim 0x41 \rightarrow 0xBE$
 - $\sim 0100\ 0001_2 \rightarrow 1011\ 1110_2$
 - $\sim 0x00 \rightarrow 0xFF$
 - $\sim 0000\ 0000_2 \rightarrow 1111\ 1111_2$
 - $0x69 \& 0x55 \rightarrow 0x41$
 - $0110\ 1001_2 \& 0101\ 0101_2 \rightarrow 0100\ 0001_2$
 - $0x69 \vee 0x55 \rightarrow 0x7D$
 - $0110\ 1001_2 \vee 0101\ 0101_2 \rightarrow 0111\ 1101_2$

Hex	Decimal	Binary
0	0	0000
1	1	0001
2	2	0010
3	3	0011
4	4	0100
5	5	0101
6	6	0110
7	7	0111
8	8	1000
9	9	1001
A	10	1010
B	11	1011
C	12	1100
D	13	1101
E	14	1110
F	15	1111

Contrast: Logic Operations in C

- Contrast to Bit-Level Operators
 - Logic Operations: $\&\&$, $\|\|$, $!$
 - View 0 as "False"
 - Anything nonzero as "True"
 - Always return 0 or 1
 - Early termination
- Examples (char data type)
 - $!0x41 \rightarrow 0x00$
 - $!0x00 \rightarrow 0x01$
 - $!0x41 \rightarrow 0x01$
 - $0x69 \&\& 0x55 \rightarrow 0x01$
 - $0x69 \|\| 0x55 \rightarrow 0x01$
 - $p \&\& *p$ (avoids null pointer access)

Watch out for $\&\&$ vs. $\&$ (and $\|\|$ vs. $\|\|$)... one of the more common oopsies in C programming

Shift Operations

- Left Shift: $x \ll y$
 - Shift bit-vector x left y positions
 - Throw away extra bits on left
 - Fill with 0's on right
- Right Shift: $x \gg y$
 - Shift bit-vector x right y positions
 - Throw away extra bits on right
 - Logical shift
 - Fill with 0's on left
 - Arithmetic shift
 - Replicate most significant bit on left
- Undefined Behavior
 - Shift amount < 0 or \geq word size

Argument x	01100010
$\ll 3$	00010000
Log. $\gg 2$	00011000
Arith. $\gg 2$	00011000

Argument x	10100010
$\ll 3$	00010000
Log. $\gg 2$	00101000
Arith. $\gg 2$	11010000

Today: Bits, Bytes, and Integers

- Representing information as bits
- Bit-level manipulations
- Integers
 - Representation: unsigned and signed
 - Conversion, casting
 - Expanding, truncating
 - Addition, negation, multiplication, shifting
 - Summary
- Representations in memory, pointers, strings
- Summary

Encoding Integers

Unsigned

$$B2U(X) = \sum_{i=0}^{w-1} x_i \cdot 2^i$$

Two's Complement

$$B2T(X) = -x_{w-1} \cdot 2^{w-1} + \sum_{i=0}^{w-2} x_i \cdot 2^i$$

```
short int x = 15213;
short int y = -15213;
```

Sign Bit

■ C does not mandate using two's complement

- But, most machines do, and we will assume so

■ C short 2 bytes long

	Decimal	Hex	Binary
x	15213	3B 6D	00111011 01101101
y	-15213	C4 93	11000100 10010011

■ Sign Bit

- For 2's complement, most significant bit indicates sign
 - 0 for nonnegative
 - 1 for negative

Two-complement: Simple Example

$$\begin{array}{ccccccc} & -16 & 8 & 4 & 2 & 1 & \\ 10 = & 0 & 1 & 0 & 1 & 0 & \end{array} \quad 8+2 = 10$$

$$\begin{array}{ccccccc} & -16 & 8 & 4 & 2 & 1 & \\ -10 = & 1 & 0 & 1 & 1 & 0 & \end{array} \quad -16+4+2 = -10$$

Two-complement Encoding Example (Cont.)

x =	15213:	00111011	01101101
y =	-15213:	11000100	10010011

Weight	15213	-15213
1	1	1
2	0	0
4	1	0
8	1	0
16	0	1
32	1	0
64	1	0
128	0	1
256	1	0
512	1	0
1024	0	1
2048	1	0
4096	1	0
8192	1	0
16384	0	1
-32768	0	1
Sum	15213	-15213

Numeric Ranges

■ Unsigned Values

- $UMin = 0$
000...0
- $UMax = 2^w - 1$
111...1

■ Two's Complement Values

- $TMin = -2^{w-1}$
100...0
- $TMax = 2^{w-1} - 1$
011...1
- Minus 1
111...1

Values for W = 16

	Decimal	Hex	Binary
UMax	65535	FF FF	11111111 11111111
TMax	32767	7F FF	01111111 11111111
TMin	-32768	80 00	10000000 00000000
-1	-1	FF FF	11111111 11111111
0	0	00 00	00000000 00000000

Values for Different Word Sizes

	8	16	32	64
UMax	255	65,535	4,294,967,295	18,446,744,073,709,551,615
TMax	127	32,767	2,147,483,647	9,223,372,036,854,775,807
TMin	-128	-32,768	-2,147,483,648	-9,223,372,036,854,775,808

■ Observations

- $|TMin| = TMax + 1$
 - Asymmetric range
- $UMax = 2 * TMax + 1$

■ C Programming

- #include <limits.h>
- Declares constants, e.g.,
 - ULONG_MAX
 - LONG_MAX
 - LONG_MIN
- Values platform specific

Unsigned & Signed Numeric Values

x	B2U(x)	B2T(x)
0000	0	0
0001	1	1
0010	2	2
0011	3	3
0100	4	4
0101	5	5
0110	6	6
0111	7	7
1000	8	-8
1001	9	-7
1010	10	-6
1011	11	-5
1100	12	-4
1101	13	-3
1110	14	-2
1111	15	-1

■ Equivalence

- Same encodings for nonnegative values

■ Uniqueness

- Every bit pattern represents unique integer value
- Each representable integer has unique bit encoding

⇒ Can Invert Mappings

- $U2B(x) = B2U^{-1}(x)$
 - Bit pattern for unsigned integer
- $T2B(x) = B2T^{-1}(x)$
 - Bit pattern for two's complement integer

Quiz Time!


Check out:

<https://canvas.cmu.edu/courses/5835>

Today: Bits, Bytes, and Integers

- Representing information as bits
- Bit-level manipulations
- Integers
 - Representation: unsigned and signed
 - Conversion, casting
 - Expanding, truncating
 - Addition, negation, multiplication, shifting
 - Summary
- Representations in memory, pointers, strings

Mapping Between Signed & Unsigned



- Mappings between unsigned and two's complement numbers: Keep bit representations and reinterpret

Mapping Signed ↔ Unsigned

Bits	Signed	Unsigned
0000	0	0
0001	1	1
0010	2	2
0011	3	3
0100	4	4
0101	5	5
0110	6	6
0111	7	7
1000	-8	8
1001	-7	9
1010	-6	10
1011	-5	11
1100	-4	12
1101	-3	13
1110	-2	14
1111	-1	15




Mapping Signed ↔ Unsigned

Bits	Signed	Unsigned
0000	0	0
0001	1	1
0010	2	2
0011	3	3
0100	4	4
0101	5	5
0110	6	6
0111	7	7
1000	-8	8
1001	-7	9
1010	-6	10
1011	-5	11
1100	-4	12
1101	-3	13
1110	-2	14
1111	-1	15

=

+/- 16


Relation between Signed & Unsigned



Conversion Visualized

■ 2's Comp. → Unsigned

- Ordering Inversion
- Negative → Big Positive



Signed vs. Unsigned in C

■ Constants

- By default are considered to be signed integers
- Unsigned if have "U" as suffix
0U, 4294967259U

■ Casting

- Explicit casting between signed & unsigned same as U2T and T2U

```
int tx, ty;
unsigned ux, uy;
tx = (int) ux;
uy = (unsigned) ty;

Implicit casting also occurs via assignments and procedure calls
tx = ux;          int fun(unsigned u);
uy = ty;          uy = fun(tx);
```

Casting Surprises

■ Expression Evaluation

- If there is a mix of unsigned and signed in single expression, *signed values implicitly cast to unsigned*
- Including comparison operations <, >, ==, <=, >=
- Examples for W = 32: TMIN = -2,147,483,648, TMAX = 2,147,483,647

Constant ₁	Constant ₂	Relation	Evaluation
0	0U	==	unsigned
-1	0	<	signed
-1	0U	>	unsigned
2147483647	-2147483647-1	>	signed
2147483647U	-2147483647-1	<	unsigned
-1	-2	>	signed
(unsigned)-1	-2	>	unsigned
2147483647	2147483648U	<	unsigned
2147483647	(int) 2147483648U	>	signed

Unsigned vs. Signed: Easy to Make Mistakes

```
unsigned i;
for (i = cnt-2; i >= 0; i--)
    a[i] += a[i+1];
```

- Can be very subtle
- ```
#define DELTA sizeof(int)
int i;
for (i = CNT; i-DELTA >= 0; i-= DELTA)
 . . .
```

## Summary

## Casting Signed ↔ Unsigned: Basic Rules

- Bit pattern is maintained
- But reinterpreted
- Can have unexpected effects: adding or subtracting 2<sup>n</sup>
- Expression containing signed and unsigned int
  - int is cast to unsigned!!


## Today: Bits, Bytes, and Integers

- Representing information as bits
- Bit-level manipulations
- Integers
  - Representation: unsigned and signed
  - Conversion, casting
  - Expanding, truncating
  - Addition, negation, multiplication, shifting
  - Summary
- Representations in memory, pointers, strings

## Sign Extension

- **Task:**
  - Given  $w$ -bit signed integer  $x$
  - Convert it to  $w+k$ -bit integer with same value


- **Rule:**
  - Make  $k$  copies of sign bit:
  - $X' = x_{w-1}, \dots, x_{w-1}, x_{w-1}, x_{w-2}, \dots, x_0$



Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

37

## Sign Extension: Simple Example



Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

38

## Larger Sign Extension Example

```
short int x = 15213;
int ix = (int) x;
short int y = -15213;
int iy = (int) y;
```

|    | Decimal | Hex         | Binary                              |
|----|---------|-------------|-------------------------------------|
| x  | 15213   | 3B 6D       | 00111011 01101101                   |
| ix | 15213   | 00 00 3B 6D | 00000000 00000000 00111011 01101101 |
| y  | -15213  | C4 93       | 11000100 10010011                   |
| iy | -15213  | FF FF C4 93 | 11111111 11111111 11000100 10010011 |

- Converting from smaller to larger integer data type
- C automatically performs sign extension


Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

39

## Truncation

- **Task:**
  - Given  $k+w$ -bit signed or unsigned integer  $X$
  - Convert it to  $w$ -bit integer  $X'$  with same value for "small enough"  $X$

- **Rule:**
  - Drop top  $k$  bits:
  - $X' = x_{w-1}, x_{w-2}, \dots, x_0$



Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

40

## Truncation: Simple Example



Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

41

## Summary:

### Expanding, Truncating: Basic Rules

- **Expanding (e.g., short int to int)**
  - Unsigned: zeros added
  - Signed: sign extension
  - Both yield expected result
- **Truncating (e.g., unsigned to unsigned short)**
  - Unsigned/signed: bits are truncated
  - Result reinterpreted
  - Unsigned: mod operation
  - Signed: similar to mod
  - For small (in magnitude) numbers yields expected behavior

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

42

## Summary of Today: Bits, Bytes, and Integers

- Representing information as bits
- Bit-level manipulations
- Integers
  - Representation: unsigned and signed
  - Conversion, casting
  - Expanding, truncating
  - Addition, negation, multiplication, shifting
- Representations in memory, pointers, strings
- Summary