



Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

1

Summary From Last Lecture

- Representing information as bits
- Bit-level manipulations
- Integers
 - Representation: unsigned and signed
 - Conversion, casting
 - Expanding, truncating
 - Addition, negation, multiplication, shifting
- Representations in memory, pointers, strings
- Summary

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

4

Sign Extension and Truncation

- Sign Extension

- Truncation

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

7



Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

2

Bits, Bytes, and Integers – Part 2

15-213: Introduction to Computer Systems
3rd Lecture, Sept. 4, 2018

Encoding Integers

Unsigned $B2U(X) = \sum_{i=0}^{w-1} x_i \cdot 2^i$	Two's Complement $B2T(X) = -x_{w-1} \cdot 2^{w-1} + \sum_{i=0}^{w-2} x_i \cdot 2^i$
--	---

Sign Bit

Two's Complement Examples (w = 5)

$-16 \quad 8 \quad 4 \quad 2 \quad 1$ $10 = 0 \quad 1 \quad 0 \quad 1 \quad 0 \quad 8+2 = 10$	$-16 \quad 8 \quad 4 \quad 2 \quad 1$ $-10 = 1 \quad 0 \quad 1 \quad 1 \quad 0 \quad -16+4+2 = -10$
--	--

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

5

Today: Bits, Bytes, and Integers

- Representing information as bits
- Bit-level manipulations
- Integers
 - Representation: unsigned and signed
 - Conversion, casting
 - Expanding, truncating
 - Addition, negation, multiplication, shifting
- Representations in memory, pointers, strings
- Summary

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

8

Assignment Announcements

- **Lab 0 available via course web page and Autolab.**
 - Due Thurs. Sept. 6, 11:59pm
 - No grace days
 - No late submissions
 - Just do it!
- **Lab 1 available via Autolab**
 - Due Thurs, Sept. 13, 11:59pm
 - Read instructions carefully: writeup, bits.c, tests.c
 - Quirky software infrastructure
 - Based on lectures 2, 3, and 4 (CS:APP Chapter 2)
 - After today's lecture you will know everything for the integer problems
 - Floating point covered Thursday Sept. 6

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

3



Carnegie Mellon

Unsigned & Signed Numeric Values

X	B2U(X)	B2T(X)
0000	0	0
0001	1	1
0010	2	2
0011	3	3
0100	4	4
0101	5	5
0110	6	6
0111	7	7
1000	8	-8
1001	9	-7
1010	10	-6
1011	11	-5
1100	12	-4
1101	13	-3
1110	14	-2
1111	15	-1

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

6

Unsigned Addition

Operands: w bits

$$\begin{array}{r} u \\ + v \\ \hline u+v \end{array}$$

True Sum: $w+1$ bits

Discard Carry: w bits $UAdd_w(u, v) = \boxed{}$

Standard Addition Function

- Ignores carry output

Implements Modular Arithmetic

$$s = UAdd_w(u, v) = u + v \bmod 2^w$$

<code>unsigned char</code>	$1110 \ 1001 \ E9 \ 223$	$\boxed{+ \ 1101 \ 0101 \ D5 \ 213}$
<code>+ 1011 1110</code>	<code>1E1110</code>	$\boxed{+ \ 1B \ 446}$
		$\boxed{= \ BE \ 190}$

Hex Decimal Binary

0	0	0000
1	1	0001
2	2	0010
3	3	0011
4	4	0100
5	5	0101
6	6	0110
7	7	0111
8	8	1000
9	9	1001
A	10	1010
B	11	1011
C	12	1100
D	13	1101
E	14	1110
F	15	1111

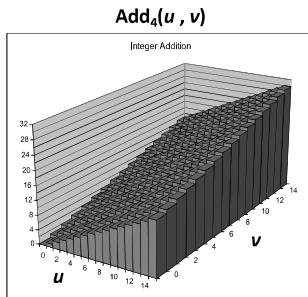
Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

9

Visualizing (Mathematical) Integer Addition

■ Integer Addition

- 4-bit integers u, v
- Compute true sum $\text{Add}_4(u, v)$
- Values increase linearly with u and v
- Forms planar surface



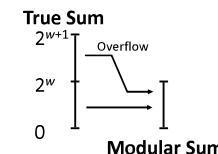
Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

10

Visualizing Unsigned Addition

■ Wraps Around

- If true sum $\geq 2^w$
- At most once



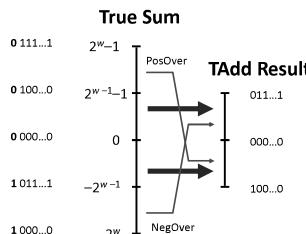
Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

11

TAdd Overflow

■ Functionality

- True sum requires $w+1$ bits
- Drop off MSB
- Treat remaining bits as 2's comp. integer



Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

13

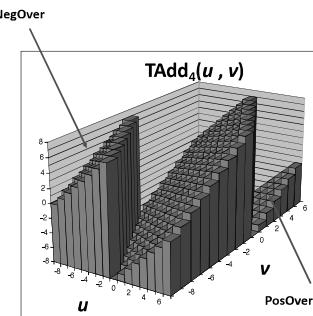
Visualizing 2's Complement Addition

■ Values

- 4-bit two's comp.
- Range from -8 to +7

■ Wraps Around

- If sum $\geq 2^{w-1}$
 - Becomes negative
 - At most once
- If sum $< -2^{w-1}$
 - Becomes positive
 - At most once



Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

14

Multiplication

■ Goal: Computing Product of w -bit numbers x, y

- Either signed or unsigned

■ But, exact results can be bigger than w bits

- Unsigned: up to $2w$ bits
 - Result range: $0 \leq x * y \leq (2^w - 1)^2 = 2^{2w} - 2^{w+1} + 1$
- Two's complement min (negative): Up to $2w-1$ bits
 - Result range: $x * y \geq (-2^{w-1})^2 (2^{w-1} - 1) = -2^{2w-2} + 2^{w-1}$
- Two's complement max (positive): Up to $2w$ bits, but only for $(TMin_w)^2$
 - Result range: $x * y \leq (-2^{w-1})^2 = 2^{2w-2}$

■ So, maintaining exact results...

- would need to keep expanding word size with each product computed
- is done in software, if needed
 - e.g., by "arbitrary precision" arithmetic packages

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

16

Unsigned Multiplication in C

Operands: w bits

$$\begin{array}{r} u \\ * v \\ \hline \text{True Product: } 2^w \text{ bits } u \cdot v \end{array} \quad \begin{array}{c} \dots \\ \dots \\ \dots \end{array}$$

Discard w bits: w bits $\text{UMult}_w(u, v) \quad \dots \quad \dots$

■ Standard Multiplication Function

- Ignores high order w bits

■ Implements Modular Arithmetic

$$\text{UMult}_w(u, v) = u \cdot v \bmod 2^w$$

$$\begin{array}{r} 1110 \ 1001 \quad E9 \quad 223 \\ * \quad 1101 \ 0101 \quad * \quad D5 \quad * \quad 213 \\ \hline 1100 \ 0001 \ 1101 \ 1101 \quad C1DD \quad 47499 \\ \hline 1101 \ 1101 \quad DD \quad 221 \end{array}$$

17

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

Two's Complement Addition

Operands: w bits

$$\begin{array}{r} u \quad \dots \quad \dots \\ + v \quad \dots \quad \dots \\ \hline u + v \end{array} \quad \begin{array}{c} \dots \\ \dots \\ \dots \end{array}$$

Discard Carry: w bits $\text{TAdd}_w(u, v) \quad \dots \quad \dots$

■ TAdd and UAdd have Identical Bit-Level Behavior

- Signed vs. unsigned addition in C:


```
int s, t, u, v;
s = (int) ((unsigned) u + (unsigned) v);
t = u + v;
```
- Will give $s == t$

$$\begin{array}{r} 1110 \ 1001 \quad E9 \quad -23 \\ + \ 1101 \ 0101 \quad + D5 \quad + -43 \\ \hline 1 \ 1011 \ 1110 \quad 1BE \quad 446 \\ 1011 \ 1110 \quad BE \quad -66 \end{array}$$

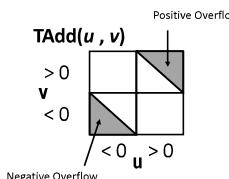
Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

12

Characterizing TAdd

■ Functionality

- True sum requires $w+1$ bits
- Drop off MSB
- Treat remaining bits as 2's comp. integer



$$\text{TAdd}_w(u, v) = \begin{cases} u + v + 2^w & u + v < TMin_w \text{ (NegOver)} \\ u + v & TMin_w \leq u + v \leq TMax_w \\ u + v - 2^w & TMax_w < u + v \text{ (PosOver)} \end{cases}$$

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

15

15

15

Signed Multiplication in C

Operands: w bits

$$\begin{array}{r} u \quad \dots \quad \dots \\ * \quad v \quad \dots \quad \dots \\ \hline u \cdot v \end{array} \quad \begin{array}{c} \dots \\ \dots \\ \dots \end{array}$$

Discard w bits: w bits $\text{TMult}_w(u, v) \quad \dots \quad \dots$

■ Standard Multiplication Function

- Ignores high order w bits
- Some of which are different for signed vs. unsigned multiplication
- Lower bits are the same

$$\begin{array}{r} 1110 \ 1001 \quad E9 \quad -23 \\ * \quad 1101 \ 0101 \quad * \quad D5 \quad * \quad -43 \\ \hline 0000 \ 0011 \ 1101 \ 1101 \quad 03DD \quad 989 \\ 1101 \ 1101 \quad DD \quad -35 \end{array}$$

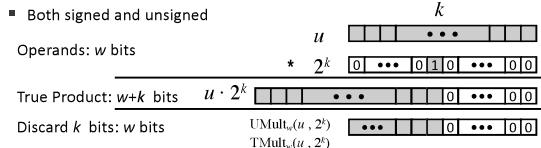
Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

18

Power-of-2 Multiply with Shift

■ Operation

- $u \ll k$ gives $u * 2^k$
- Both signed and unsigned



■ Examples

- $u \ll 3 == u * 8$
- $(u \ll 5) - (u \ll 3) == u * 24$
- Most machines shift and add faster than multiply
- Compiler generates this code automatically

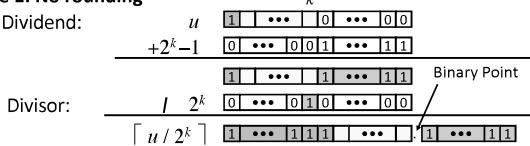
Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

Correct Power-of-2 Divide

■ Quotient of Negative Number by Power of 2

- Want $\lceil x / 2^k \rceil$ (Round Toward 0)
- Compute as $\lfloor (x+2^k-1) / 2^k \rfloor$
 - In C: $(x + (1 < k) - 1) \gg k$
 - Biases dividend toward 0

Case 1: No rounding

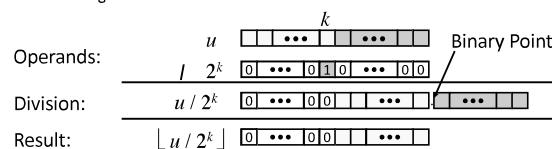
*Biasing has no effect*

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

Unsigned Power-of-2 Divide with Shift

■ Quotient of Unsigned by Power of 2

- $u \gg k$ gives $\lfloor u / 2^k \rfloor$
- Uses logical shift

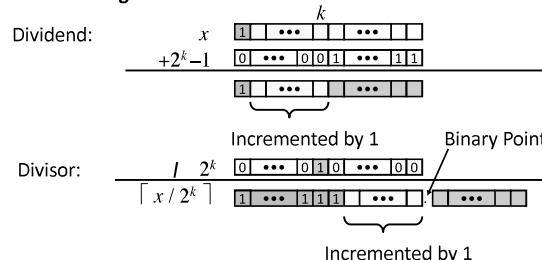


	Division	Computed	Hex	Binary
x	15213	15213	3B 6D	00111011 01101101
$x \gg 1$	7606.5	7606	1D B6	00011101 10110110
$x \gg 4$	950.8125	950	03 B6	00000011 10110110
$x \gg 8$	59.4257813	59	00 3B	00000000 00111011

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

Correct Power-of-2 Divide (Cont.)

Case 2: Rounding

*Incremented by 1 adds 1 to final result*

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

Complement & Increment Examples

 $x = 0$

	Decimal	Hex	Binary
0	0	00 00	00000000 00000000
~ 0	-1	FF FF	11111111 11111111
$\sim 0 + 1$	0	00 00	00000000 00000000

 $x = TMin$

	Decimal	Hex	Binary
x	-32768	80 00	10000000 00000000
$\sim x$	32767	7F FF	01111111 11111111
$\sim x + 1$	-32768	80 00	10000000 00000000

Canonical counter example

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

Today: Bits, Bytes, and Integers

- Representing information as bits
- Bit-level manipulations
- Integers**
 - Representation: unsigned and signed
 - Conversion, casting
 - Expanding, truncating
 - Addition, negation, multiplication, shifting
 - Summary**
- Representations in memory, pointers, strings

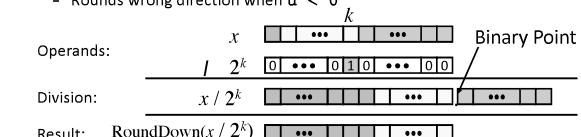
Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

Signed Power-of-2 Divide with Shift

■ Quotient of Signed by Power of 2

- $x \gg k$ gives $\lfloor x / 2^k \rfloor$

Uses arithmetic shift

Rounds wrong direction when $x < 0$ 

	Division	Computed	Hex	Binary
y	-15213	-15213	C4 93	11000100 10010011
$y \gg 1$	-7606.5	-7607	E2 49	11100010 01001001
$y \gg 4$	-950.8125	-951	FC 49	11111100 01001001
$y \gg 8$	-59.4257813	-60	FF C4	11111111 11000100

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

Negation: Complement & Increment

■ Negate through complement and increase

$$\sim x + 1 = -x$$

■ Example

- Observation: $\sim x + x == 1111...111 == -1$

$$\begin{array}{r} x \quad 100111101 \\ + \sim x \quad 01100010 \\ \hline -1 \quad 111111111 \end{array}$$

 $x = 15213$

	Decimal	Hex	Binary
x	15213	3B 6D	00111011 01101101
$\sim x$	-15214	C4 92	11000100 10010010
$\sim x + 1$	-15213	C4 93	11000100 10010011
y	-15213	C4 93	11000100 10010011

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

Arithmetic: Basic Rules

■ Addition:

- Unsigned/signed: Normal addition followed by truncate, same operation on bit level
- Unsigned: addition mod 2^w
 - Mathematical addition + possible subtraction of 2^w
- Signed: modified addition mod 2^w (result in proper range)
 - Mathematical addition + possible addition or subtraction of 2^w

■ Multiplication:

- Unsigned/signed: Normal multiplication followed by truncate, same operation on bit level
- Unsigned: multiplication mod 2^w
- Signed: modified multiplication mod 2^w (result in proper range)

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

Why Should I Use Unsigned?

- **Don't use without understanding implications**

- Easy to make mistakes


```
unsigned i;
for (i = cnt-2; i >= 0; i--)
    a[i] += a[i+1];
```

- Can be very subtle


```
#define DELTA sizeof(int)
int i;
for (i = CNT; i-DELTA >= 0; i-= DELTA)
    ...
```

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

28

Quiz Time!

Check out:

<https://canvas.cmu.edu/courses/5835>

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

31

Machine Words

- Any given computer has a "Word Size"

- Nominal size of integer-valued data
 - and of addresses
- Until recently, most machines used 32 bits (4 bytes) as word size
 - Limits addresses to 4GB (2^{32} bytes)
- Increasingly, machines have 64-bit word size
 - Potentially, could have 18 EB (exabytes) of addressable memory
 - That's 18.4×10^{18}
- Machines still support multiple data formats
 - Fractions or multiples of word size
 - Always integral number of bytes

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

34

Counting Down with Unsigned

- Proper way to use unsigned as loop index

```
unsigned i;
for (i = cnt-2; i < cnt; i--)
    a[i] += a[i+1];
```

- See Robert Seacord, *Secure Coding in C and C++*

- C Standard guarantees that unsigned addition will behave like modular arithmetic
 - $0 - 1 \rightarrow UMax$

- Even better

```
size_t i;
for (i = cnt-2; i < cnt; i--)
    a[i] += a[i+1];
■ Data type size_t defined as unsigned value with length = word size
■ Code will work even if cnt = UMax
■ What if cnt is signed and < 0?
```

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

29

Today: Bits, Bytes, and Integers

- Representing information as bits
- Bit-level manipulations
- Integers
 - Representation: unsigned and signed
 - Conversion, casting
 - Expanding, truncating
 - Addition, negation, multiplication, shifting
 - Summary
- Representations in memory, pointers, strings

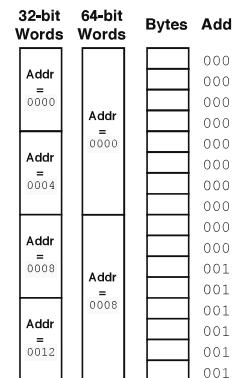
Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

32

Word-Oriented Memory Organization

- Addresses Specify Byte Locations

- Address of first byte in word
- Addresses of successive words differ by 4 (32-bit) or 8 (64-bit)



35

Why Should I Use Unsigned? (cont.)

- **Do Use When Performing Modular Arithmetic**
 - Multiprecision arithmetic

- **Do Use When Using Bits to Represent Sets**
 - Logical right shift, no sign extension

- **Do Use In System Programming**
 - Bit masks, device commands,...

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

30

Byte-Oriented Memory Organization



- **Programs refer to data by address**

- Conceptually, envision it as a very large array of bytes
 - In reality, it's not, but can think of it that way
- An address is like an index into that array
 - and, a pointer variable stores an address

- **Note: system provides private address spaces to each "process"**

- Think of a process as a program being executed
- So, a program can clobber its own data, but not that of others

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

33

Example Data Representations

C Data Type	Typical 32-bit	Typical 64-bit	x86-64
<code>char</code>	1	1	1
<code>short</code>	2	2	2
<code>int</code>	4	4	4
<code>long</code>	4	8	8
<code>float</code>	4	4	4
<code>double</code>	8	8	8
<code>pointer</code>	4	8	8

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

36

Byte Ordering

- So, how are the bytes within a multi-byte word ordered in memory?
- Conventions
 - Big Endian: Sun (Oracle SPARC), PPC Mac, Internet
 - Least significant byte has highest address
 - Little Endian: x86, ARM processors running Android, iOS, and Linux
 - Least significant byte has lowest address

Examining Data Representations

- Code to Print Byte Representation of Data
 - Casting pointer to unsigned char * allows treatment as a byte array

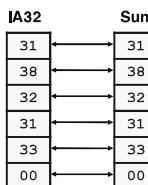
```
typedef unsigned char *pointer;

void show_bytes(pointer start, size_t len){
    size_t i;
    for (i = 0; i < len; i++)
        printf("%p\t%02x\n", start+i, start[i]);
    printf("\n");
}
```

Printf directives:
`%p`: Print pointer
`%x`: Print Hexadecimal

Representing Strings

- Strings in C
 - Represented by array of characters
 - Each character encoded in ASCII format
 - Standard 7-bit encoding of character set
 - Character "0" has code 0x30
 - Digit *i* has code 0x30+i
 - String should be null-terminated
 - Final character = 0
- Compatibility
 - Byte ordering not an issue



Byte Ordering Example

Example

- Variable *x* has 4-byte value of 0x01234567
- Address given by `&x` is 0x100

Big Endian

	0x100	0x101	0x102	0x103	
	01	23	45	67	

Little Endian

	0x100	0x101	0x102	0x103	
	67	45	23	01	

show_bytes Execution Example

```
int a = 15213;
printf("int a = 15213;\n");
show_bytes((pointer) &a, sizeof(int));
```

Result (Linux x86-64):

```
int a = 15213;
0x7fff7f71dbc 6d
0x7fff7f71dbd 3b
0x7fff7f71dbe 00
0x7fff7f71dbf 00
```

Reading Byte-Reversed Listings

Disassembly

- Text representation of binary machine code
- Generated by program that reads the machine code

Example Fragment

Address	Instruction Code	Assembly Rendition
8048365:	5b	pop %ebx
8048366:	81 c3 ab 12 00 00	add \$0x12ab,%ebx
804836c:	83 bb 28 00 00 00 00	cmpl \$0x0,0x28(%ebx)

Dic平ering Numbers

- Value:
- Pad to 32 bits:
- Split into bytes:
- Reverse:

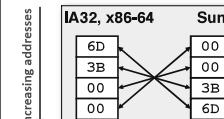
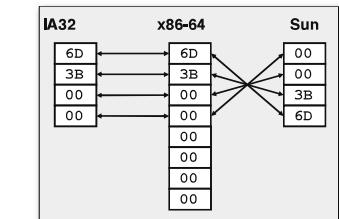
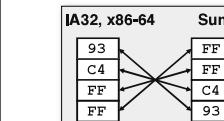
```
0x12ab
0x000012ab
00 00 12 ab
ab 12 00 00
```

Representing Integers

Decimal: 15213

Binary: 0011 1011 0110 1101

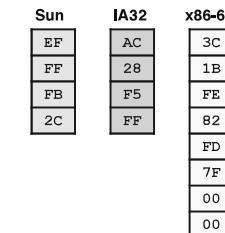
Hex: 3 B 6 D

`int A = 15213;``long int C = 15213;``int B = -15213;`

Two's complement representation

Representing Pointers

```
int B = -15213;
int *P = &B;
```



Different compilers & machines assign different locations to objects

Even get different results each time run program

Integer C Puzzles

<code>x < 0</code>	$\Rightarrow ((x*2) < 0)$	
<code>ux >= 0</code>	$\Rightarrow (ux >= 0)$	
<code>x & 7 == 7</code>	$\Rightarrow (x << 30) < 0$	
<code>ux > -1</code>	$\Rightarrow (ux > -1)$	
<code>x > y</code>	$\Rightarrow -x < -y$	
<code>x * x >= 0</code>	$\Rightarrow x * x >= 0$	
<code>x > 0 && y > 0</code>	$\Rightarrow x + y > 0$	
<code>x >= 0</code>	$\Rightarrow -x <= 0$	
<code>x <= 0</code>	$\Rightarrow -x >= 0$	
<code>(x -x)>>31 == -1</code>	$\Rightarrow (x -x)>>31 == -1$	
<code>ux >> 3 == ux/8</code>	$\Rightarrow ux >> 3 == ux/8$	
<code>x >> 3 == x/8</code>	$\Rightarrow x >> 3 == x/8$	
<code>x & (x-1) != 0</code>	$\Rightarrow x & (x-1) != 0$	

Summary

- Representing information as bits
- Bit-level manipulations
- Integers
 - Representation: unsigned and signed
 - Conversion, casting
 - Expanding, truncating
 - Addition, negation, multiplication, shifting
- Representations in memory, pointers, strings
- Summary