



Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

---

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition 3

---

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition 4

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition 7

---

Bryant and O'Hallaron, *Computer Systems: A Programmer's Perspective*, Third Edition

---

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition 9

## Intel's 64-Bit History

- **2001: Intel Attempts Radical Shift from IA32 to IA64**
  - Totally different architecture (Itanium)
  - Executes IA32 code only as legacy
  - Performance disappointing
- **2003: AMD Steps in with Evolutionary Solution**
  - x86-64 (now called "AMD64")
- **Intel Felt Obligated to Focus on IA64**
  - Hard to admit mistake or that AMD is better
- **2004: Intel Announces EM64T extension to IA32**
  - Extended Memory 64-bit Technology
  - Almost identical to x86-64!
- **All but low-end x86 processors support x86-64**
  - But, lots of code still runs in 32-bit mode

Reiser and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

## Our Coverage

- **IA32**
  - The traditional x86
  - For 15/18-213: RIP, Summer 2015
- **x86-64**
  - The standard
  - `shark> gcc hello.c`
  - `shark> gcc -m64 hello.c`
- **Presentation**
  - Book covers x86-64
  - Web aside on IA32
  - We will only cover x86-64

Reiser and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

## Today: Machine Programming I: Basics

- History of Intel processors and architectures
- **Assembly Basics: Registers, operands, move**
- Arithmetic & logical operations
- C, assembly, machine code

Reiser and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

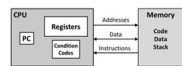
## Levels of Abstraction

C programmer



Nice clean layers,  
but beware...

Assembly programmer



Computer Designer

Caches, clock freq, layout, ...

Of course, you know that: It's why you are taking this course.

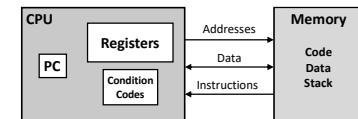
Reiser and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

## Definitions

- **Architecture: (also ISA: instruction set architecture)** The parts of a processor design that one needs to understand for writing correct machine/assembly code
  - Examples: instruction set specification, registers
- **Machine Code:** The byte-level programs that a processor executes
- **Assembly Code:** A text representation of machine code
- **Microarchitecture: Implementation of the architecture**
  - Examples: cache sizes and core frequency
- **Example ISAs:**
  - Intel: x86, IA32, Itanium, x86-64
  - ARM: Used in almost all mobile phones
  - RISC V: New open-source ISA

Reiser and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

## Assembly/Machine Code View



### Programmer-Visible State

- **PC: Program counter**
  - Address of next instruction
  - Called "RIP" (x86-64)
- **Register file**
  - Heavily used program data
- **Condition codes**
  - Store status information about most recent arithmetic or logical operation
  - Used for conditional branching
- **Memory**
  - Byte addressable array
  - Code and user data
  - Stack to support procedures

Reiser and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

## Assembly Characteristics: Data Types

- **"Integer" data of 1, 2, 4, or 8 bytes**
  - Data values
  - Addresses (untyped pointers)
- Floating point data of 4, 8, or 10 bytes
- (SIMD vector data types of 8, 16, 32 or 64 bytes)
- **Code: Byte sequences encoding series of instructions**
- **No aggregate types such as arrays or structures**
  - Just contiguously allocated bytes in memory

Reiser and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

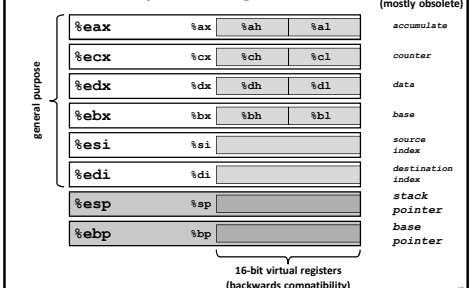
## x86-64 Integer Registers

<code>%rax</code>	<code>%eax</code>	<code>%r8</code>	<code>%r8d</code>
<code>%rbx</code>	<code>%ebx</code>	<code>%r9</code>	<code>%r9d</code>
<code>%rcx</code>	<code>%ecx</code>	<code>%r10</code>	<code>%r10d</code>
<code>%rdx</code>	<code>%edx</code>	<code>%r11</code>	<code>%r11d</code>
<code>%rsi</code>	<code>%esi</code>	<code>%r12</code>	<code>%r12d</code>
<code>%rdi</code>	<code>%edi</code>	<code>%r13</code>	<code>%r13d</code>
<code>%rsp</code>	<code>%esp</code>	<code>%r14</code>	<code>%r14d</code>
<code>%rbp</code>	<code>%ebp</code>	<code>%r15</code>	<code>%r15d</code>

- Can reference low-order 4 bytes (also low-order 1 & 2 bytes)
- Not part of memory (or cache)

Reiser and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

## Some History: IA32 Registers



Reiser and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

## Assembly Characteristics: Operations

- Transfer data between memory and register
  - Load data from memory into register
  - Store register data into memory
- Perform arithmetic function on register or memory data
- Transfer control
  - Unconditional jumps to/from procedures
  - Conditional branches
  - Indirect branches

Repart and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

19

## Moving Data

- Moving Data
  - movq source, Dest
- Operand Types
  - Immediate: Constant integer data
    - Example: \$0x400, \$-533
    - Like C constant, but prefixed with '\$'
    - Encoded with 1, 2, or 4 bytes
  - Register: One of 16 integer registers
    - Example: %rax, %r13
    - But %rsp reserved for special use
    - Others have special uses for particular instructions
  - Memory: Consecutive bytes of memory at address given by register
    - Simplest example: (%rax)
    - Various other "addressing modes"

%rax
%rcx
%rdx
%rbx
%rsi
%rdi
%rsp
%rbp
%rN

Warning: Intel does use  
mov Dest, Source

Repart and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

20

## movq Operand Combinations

	Source	Dest	Src, Dest	C Analog
movq	Imm	Reg	movq \$0x4, %rax	temp = 0x4;
		Mem	movq \$-147, (%rax)	*p = -147;
	Reg	Reg	movq %rax, %rdx	temp2 = temp1;
		Mem	movq %rax, (%rdx)	*p = temp;
	Mem	Reg	movq (%rax), %rdx	temp = *p;

Cannot do memory-memory transfer with a single instruction

Repart and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

21

## Simple Memory Addressing Modes

- Normal (R) Mem[Reg[R]]
  - Register R specifies memory address
  - Aha! Pointer dereferencing in C

```
movq (%rcx), %rax
```
- Displacement D(R) Mem[Reg[R]+D]
  - Register R specifies start of memory region
  - Constant displacement D specifies offset

```
movq 8(%rbp), %rdx
```

Repart and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

22

## Example of Simple Addressing Modes

```
void
whatAmI(<type> a, <type> b)
{
    ???
}
```

whatAmI:

```
movq    (%rdi), %rax
movq    (%rsi), %rdx
movq    %rdx, (%rdi)
movq    %rax, (%rsi)
ret
```

Diagram showing %rdi pointing to 'a' and %rsi pointing to 'b'.

Repart and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

23

## Example of Simple Addressing Modes

```
void swap
(long *xp, long *yp)
{
    long t0 = *xp;
    long t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```

swap:

```
movq    (%rdi), %rax
movq    (%rsi), %rdx
movq    %rdx, (%rdi)
movq    %rax, (%rsi)
ret
```

Repart and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

24

## Understanding Swap()

```
void swap
(long *xp, long *yp)
{
    long t0 = *xp;
    long t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```

Registers

%rdi	xp
%rsi	yp
%rax	t0
%rdx	t1

Memory

123
456

swap:

```
movq    (%rdi), %rax # t0 = *xp
movq    (%rsi), %rdx # t1 = *yp
movq    %rdx, (%rdi) # *xp = t1
movq    %rax, (%rsi) # *yp = t0
ret
```

Repart and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

25

## Understanding Swap()

Registers

%rdi	0x120
%rsi	0x100
%rax	
%rdx	

Memory

123	0x120
	0x118
	0x110
	0x108
456	0x100

swap:

```
movq    (%rdi), %rax # t0 = *xp
movq    (%rsi), %rdx # t1 = *yp
movq    %rdx, (%rdi) # *xp = t1
movq    %rax, (%rsi) # *yp = t0
ret
```

Repart and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

26

## Understanding Swap()

Registers

%rdi	0x120
%rsi	0x100
%rax	123
%rdx	

Memory

123	0x120
	0x118
	0x110
	0x108
456	0x100

swap:

```
movq    (%rdi), %rax # t0 = *xp
movq    (%rsi), %rdx # t1 = *yp
movq    %rdx, (%rdi) # *xp = t1
movq    %rax, (%rsi) # *yp = t0
ret
```

Repart and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

27

## Understanding Swap()

Registers		Memory	
	Address		Address
%rdi	0x120	123	0x120
%rsi	0x100		0x118
%rax	123		0x110
%rdx	456	456	0x108
			0x100

```

swap:
    movq    (%rdi), %rax    # t0 = *xp
    movq    (%rsi), %rdx    # t1 = *yp
    movq    %rdx, (%rdi)    # *xp = t1
    movq    %rax, (%rsi)    # *yp = t0
    ret
    
```

Revised and O'Reilly, Computer Systems: A Programmer's Perspective, Third Edition

21

## Understanding Swap()

Registers		Memory	
	Address		Address
%rdi	0x120	456	0x120
%rsi	0x100		0x118
%rax	123		0x110
%rdx	456	456	0x108
			0x100

```

swap:
    movq    (%rdi), %rax    # t0 = *xp
    movq    (%rsi), %rdx    # t1 = *yp
    movq    %rdx, (%rdi)    # *xp = t1
    movq    %rax, (%rsi)    # *yp = t0
    ret
    
```

Revised and O'Reilly, Computer Systems: A Programmer's Perspective, Third Edition

22

## Understanding Swap()

Registers		Memory	
	Address		Address
%rdi	0x120	456	0x120
%rsi	0x100		0x118
%rax	123		0x110
%rdx	456	123	0x108
			0x100

```

swap:
    movq    (%rdi), %rax    # t0 = *xp
    movq    (%rsi), %rdx    # t1 = *yp
    movq    %rdx, (%rdi)    # *xp = t1
    movq    %rax, (%rsi)    # *yp = t0
    ret
    
```

Revised and O'Reilly, Computer Systems: A Programmer's Perspective, Third Edition

23

## Simple Memory Addressing Modes

- Normal (R) Mem[Reg[R]]
  - Register R specifies memory address
  - Aha! Pointer dereferencing in C

```
movq (%rcx), %rax
```

- Displacement D(R) Mem[Reg[R]+D]
  - Register R specifies start of memory region
  - Constant displacement D specifies offset

```
movq 8(%rbp), %rdx
```

Revised and O'Reilly, Computer Systems: A Programmer's Perspective, Third Edition

24

## Complete Memory Addressing Modes

### Most General Form

D(Rb,Ri,S) Mem[Reg[Rb]+S\*Reg[Ri]+D]

- D: Constant "displacement" 1, 2, or 4 bytes
- Rb: Base register: Any of 16 integer registers
- Ri: Index register: Any, except for %rsp
- S: Scale: 1, 2, 4, or 8 (why these numbers?)

### Special Cases

(Rb,Ri) Mem[Reg[Rb]+Reg[Ri]]  
 D(Rb,Ri) Mem[Reg[Rb]+Reg[Ri]+D]  
 (Rb,Ri,S) Mem[Reg[Rb]+S\*Reg[Ri]]

Revised and O'Reilly, Computer Systems: A Programmer's Perspective, Third Edition

25

## Address Computation Examples

%rdx	0xf000
%rcx	0x0100

- D(Rb,Ri,S) Mem[Reg[Rb]+S\*Reg[Ri]+D]
- D: Constant "displacement" 1, 2, or 4 bytes
  - Rb: Base register: Any of 16 integer registers
  - Ri: Index register: Any, except for %rsp
  - S: Scale: 1, 2, 4, or 8 (why these numbers?)

Expression	Address Computation	Address
0x8(%rdx)		
(%rdx,%rcx)		
(%rdx,%rcx,4)		
0x80(,%rdx,2)		

Revised and O'Reilly, Computer Systems: A Programmer's Perspective, Third Edition

26

## Address Computation Examples

%rdx	0xf000
%rcx	0x0100

Expression	Address Computation	Address
0x8(%rdx)	0xf000 + 0x8	0xf008
(%rdx,%rcx)	0xf000 + 0x100	0xf100
(%rdx,%rcx,4)	0xf000 + 4*0x100	0xf400
0x80(,%rdx,2)	2*0xf000 + 0x80	0x1e080

Revised and O'Reilly, Computer Systems: A Programmer's Perspective, Third Edition

27

## Today: Machine Programming I: Basics

- History of Intel processors and architectures
- Assembly Basics: Registers, operands, move
- Arithmetic & logical operations
- C, assembly, machine code

Revised and O'Reilly, Computer Systems: A Programmer's Perspective, Third Edition

28

## Address Computation Instruction

- leaq Src, Dst
  - Src is address mode expression
  - Set Dst to address denoted by expression
- Uses
  - Computing addresses without a memory reference
    - E.g., translation of  $p = 4x[i]$ ;
  - Computing arithmetic expressions of the form  $x + k*y$ 
    - $k = 1, 2, 4, \text{ or } 8$

### Example

```

long m12(long x)
{
    return x*12;
}
    
```

Converted to ASM by compiler:

```

leaq (%rdi,%rdi,2), %rax    # t = x*2*x
sllq $2, %rax               # return t<<2
    
```

Revised and O'Reilly, Computer Systems: A Programmer's Perspective, Third Edition

29

## Some Arithmetic Operations

### Two Operand Instructions:

Format	Computation	
addq <i>Src, Dest</i>	Dest = Dest + Src	
subq <i>Src, Dest</i>	Dest = Dest - Src	
imulq <i>Src, Dest</i>	Dest = Dest * Src	
salq <i>Src, Dest</i>	Dest = Dest << Src	Also called shlq
sarq <i>Src, Dest</i>	Dest = Dest >> Src	Arithmetic
shrq <i>Src, Dest</i>	Dest = Dest >> Src	Logical
xorq <i>Src, Dest</i>	Dest = Dest ^ Src	
andq <i>Src, Dest</i>	Dest = Dest & Src	
orq <i>Src, Dest</i>	Dest = Dest   Src	

### Watch out for argument order! *Src, Dest* (Warning: Intel docs use "op *Dest, Src*")

### No distinction between signed and unsigned int (why?)

## Quiz Time!

Check out:

<https://canvas.cmu.edu/courses/5835>

## Some Arithmetic Operations

### One Operand Instructions

incq <i>Dest</i>	Dest = Dest + 1
decq <i>Dest</i>	Dest = Dest - 1
negq <i>Dest</i>	Dest = -Dest
notq <i>Dest</i>	Dest = ~Dest

### See book for more instructions

## Arithmetic Expression Example

```
long arith
(long x, long y, long z)
{
    long t1 = x+y;
    long t2 = z+t1;
    long t3 = x+4;
    long t4 = y * 48;
    long t5 = t3 + t4;
    long rval = t2 * t5;
    return rval;
}
```

```
arith:
    leaq    (%rdi,%rsi), %rax
    addq    %rdx, %rax
    leaq    (%rsi,%rsi,2), %rdx
    salq    $4, %rdx
    leaq    4(%rdi,%rdx), %rcx
    imulq    %rcx, %rax
    ret
```

### Interesting Instructions

- leaq: address computation
- salq: shift
  - But, only used once
- imulq: multiplication

## Understanding Arithmetic Expression Example

```
long arith
(long x, long y, long z)
{
    long t1 = x+y;
    long t2 = z+t1;
    long t3 = x+4;
    long t4 = y * 48;
    long t5 = t3 + t4;
    long rval = t2 * t5;
    return rval;
}
```

```
arith:
    leaq    (%rdi,%rsi), %rax # t1
    addq    %rdx, %rax        # t2
    leaq    (%rsi,%rsi,2), %rdx # t4
    salq    $4, %rdx          # t5
    leaq    4(%rdi,%rdx), %rcx # t5
    imulq    %rcx, %rax        # rval
    ret
```

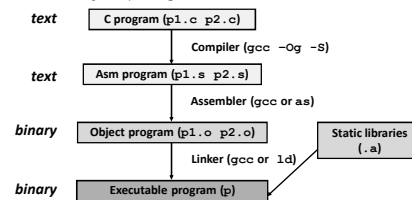
Register	Use(s)
%rdi	Argument x
%rsi	Argument y
%rdx	Argument z, t4
%rax	t1, t2, rval
%rcx	t5

## Today: Machine Programming I: Basics

- History of Intel processors and architectures
- Assembly Basics: Registers, operands, move
- Arithmetic & logical operations
- C, assembly, machine code

## Turning C into Object Code

- Code in files `p1.c p2.c`
- Compile with command: `gcc -Og p1.c p2.c -o p`
  - Use basic optimizations (`-Og`) [New to recent versions of GCC]
  - Put resulting binary in file `p`



## Compiling Into Assembly

### C Code (sum.c)

```
long plus(long x, long y);

void sumstore(long x, long y,
              long *dest)
{
    long t = plus(x, y);
    *dest = t;
}
```

### Generated x86-64 Assembly

```
sumstore:
    pushq   %rbx
    movq    %rdx, %rbx
    call    plus
    movq    %rax, (%rbx)
    popq    %rbx
    ret
```

Obtain (on shark machine) with command

`gcc -Og -S sum.c`

Produces file `sum.s`

**Warning:** Will get very different results on non-Shark machines (Andrew Linux, Mac OS-X, ...) due to different versions of gcc and different compiler settings.

## What it really looks like

```

.globl sumstore
.type sumstore, @function

sumstore:
.LFB35:
    .cfi_startproc
    pushq   %rbx
    .cfi_def_cfa_offset 16
    .cfi_offset 3, -16
    movq    %rdx, %rbx
    call    plus
    movq    %rax, (%rbx)
    popq    %rbx
    .cfi_def_cfa_offset 8
    ret
    .cfi_endproc
.LFE35:
    .size sumstore, .-sumstore

```

## What it really looks like

```
.globl sumstore
.type sumstore, @function
sumstore:
.LFB35:
.cfi_startproc
pushq %rbx
.cfi_def_cfa_offset 16
.cfi_offset 3, -16
movq %rdx, %rbx
call plus
movq %rax, (%rbx)
popq %rbx
.cfi_def_cfa_offset 8
ret
.cfi_endproc
.LFE35:
.size sumstore, .-sumstore
```

Things that look weird  
and are preceded by a '.'  
are generally directives.

```
sumstore:
pushq %rbx
movq %rdx, %rbx
call plus
movq %rax, (%rbx)
popq %rbx
ret
```

## Assembly Characteristics: Data Types

- "Integer" data of 1, 2, 4, or 8 bytes
  - Data values
  - Addresses (untyped pointers)
- Floating point data of 4, 8, or 10 bytes
- (SIMD) vector data types of 8, 16, 32 or 64 bytes
- Code: Byte sequences encoding series of instructions
- No aggregate types such as arrays or structures
  - Just contiguously allocated bytes in memory

## Assembly Characteristics: Operations

- Transfer data between memory and register
  - Load data from memory into register
  - Store register data into memory
- Perform arithmetic function on register or memory data
- Transfer control
  - Unconditional jumps to/from procedures
  - Conditional branches

## Object Code

### Code for sumstore

```
0x0400595:
0x53
0x48
0x89
0xd3
0xe8
0xf2
0xff
0xff
0xff
0x48
0x89
0x03
0x5b
0xc3
```

- Assembler
  - Translates .s into .o
  - Binary encoding of each instruction
  - Nearly-complete image of executable code
  - Missing linkages between code in different files
- Linker
  - Resolves references between files
  - Combines with static run-time libraries
    - E.g., code for malloc, printf
  - Some libraries are *dynamically linked*
  - Linking occurs when program begins execution

## Machine Instruction Example

```
*dest = t;
```

```
movq %rax, (%rbx)
```

```
0x40059e: 48 89 03
```

- C Code
  - Store value t where designated by dest
- Assembly
  - Move 8-byte value to memory
    - Quad words in x86-64 parlance
  - Operands:
    - t: Register %rax
    - dest: Register %rbx
    - \*dest: Memory M[%rbx]
- Object Code
  - 3-byte instruction
  - Stored at address 0x40059e

## Disassembling Object Code

### Disassembled

```
000000000400595 <sumstore>:
400595: 53          push %rbx
400596: 48 89 d3    mov %rdx,%rbx
400599: e8 f2 ff ff callq 400590 <plus>
40059e: 48 89 03    mov %rax, (%rbx)
4005a1: 5b         pop %rbx
4005a2: c3         retq
```

- Disassembler
  - objdump -d sum
  - Useful tool for examining object code
  - Analyzes bit pattern of series of instructions
  - Produces approximate rendition of assembly code
  - Can be run on either a.out (complete executable) or .o file

## Alternate Disassembly

### Disassembled

```
Dump of assembler code for function sumstore:
0x000000000400595 <+0>: push %rbx
0x000000000400596 <+1>: mov %rdx,%rbx
0x000000000400599 <+4>: callq 0x400590 <plus>
0x00000000040059e <+9>: mov %rax, (%rbx)
0x0000000004005a1 <+12>: pop %rbx
0x0000000004005a2 <+13>: retq
```

- Within gdb Debugger
    - Disassemble procedure
- ```
gdb sum
disassemble sumstore
```

## Alternate Disassembly

### Disassembled

#### Object Code

```
0x0400595:
0x53
0x48
0x89
0xd3
0xe8
0xf2
0xff
0xff
0xff
0x48
0x89
0x03
0x5b
0xc3
```

```
Dump of assembler code for function sumstore:
0x000000000400595 <+0>: push %rbx
0x000000000400596 <+1>: mov %rdx,%rbx
0x000000000400599 <+4>: callq 0x400590 <plus>
0x00000000040059e <+9>: mov %rax, (%rbx)
0x0000000004005a1 <+12>: pop %rbx
0x0000000004005a2 <+13>: retq
```

- Within gdb Debugger
    - Disassemble procedure
- ```
gdb sum
disassemble sumstore
x/14xb sumstore
```

## What Can be Disassembled?

```
% objdump -d WINWORD.EXE
WINWORD.EXE: file format pei-i386
```

```
No symbols in "WINWORD.EXE".
Disassembly of section .text:
```

```
30001000 <.text>:
30001000:
30001001:
30001003:
30001005:
3000100a:
```

Reverse engineering forbidden by  
Microsoft End User License Agreement

- Anything that can be interpreted as executable code
- Disassembler examines bytes and reconstructs assembly source

## Machine Programming I: Summary

- **History of Intel processors and architectures**
  - Evolutionary design leads to many quirks and artifacts
- **C, assembly, machine code**
  - New forms of visible state: program counter, registers, ...
  - Compiler must transform statements, expressions, procedures into low-level instruction sequences
- **Assembly Basics: Registers, operands, move**
  - The x86-64 move instructions cover wide range of data movement forms
- **Arithmetic**
  - C compiler will figure out different instruction combinations to carry out computation